Performance
at MongoDB

Understanding and Improving Performance at MongoDB

David Daly
Staff Performance Engineer
https://daviddaly.me/

Agenda

Basics of Performance Testing

Scaling Our Testing

Adding Features + Updates

Community Interactions

Basics of Performance Testing

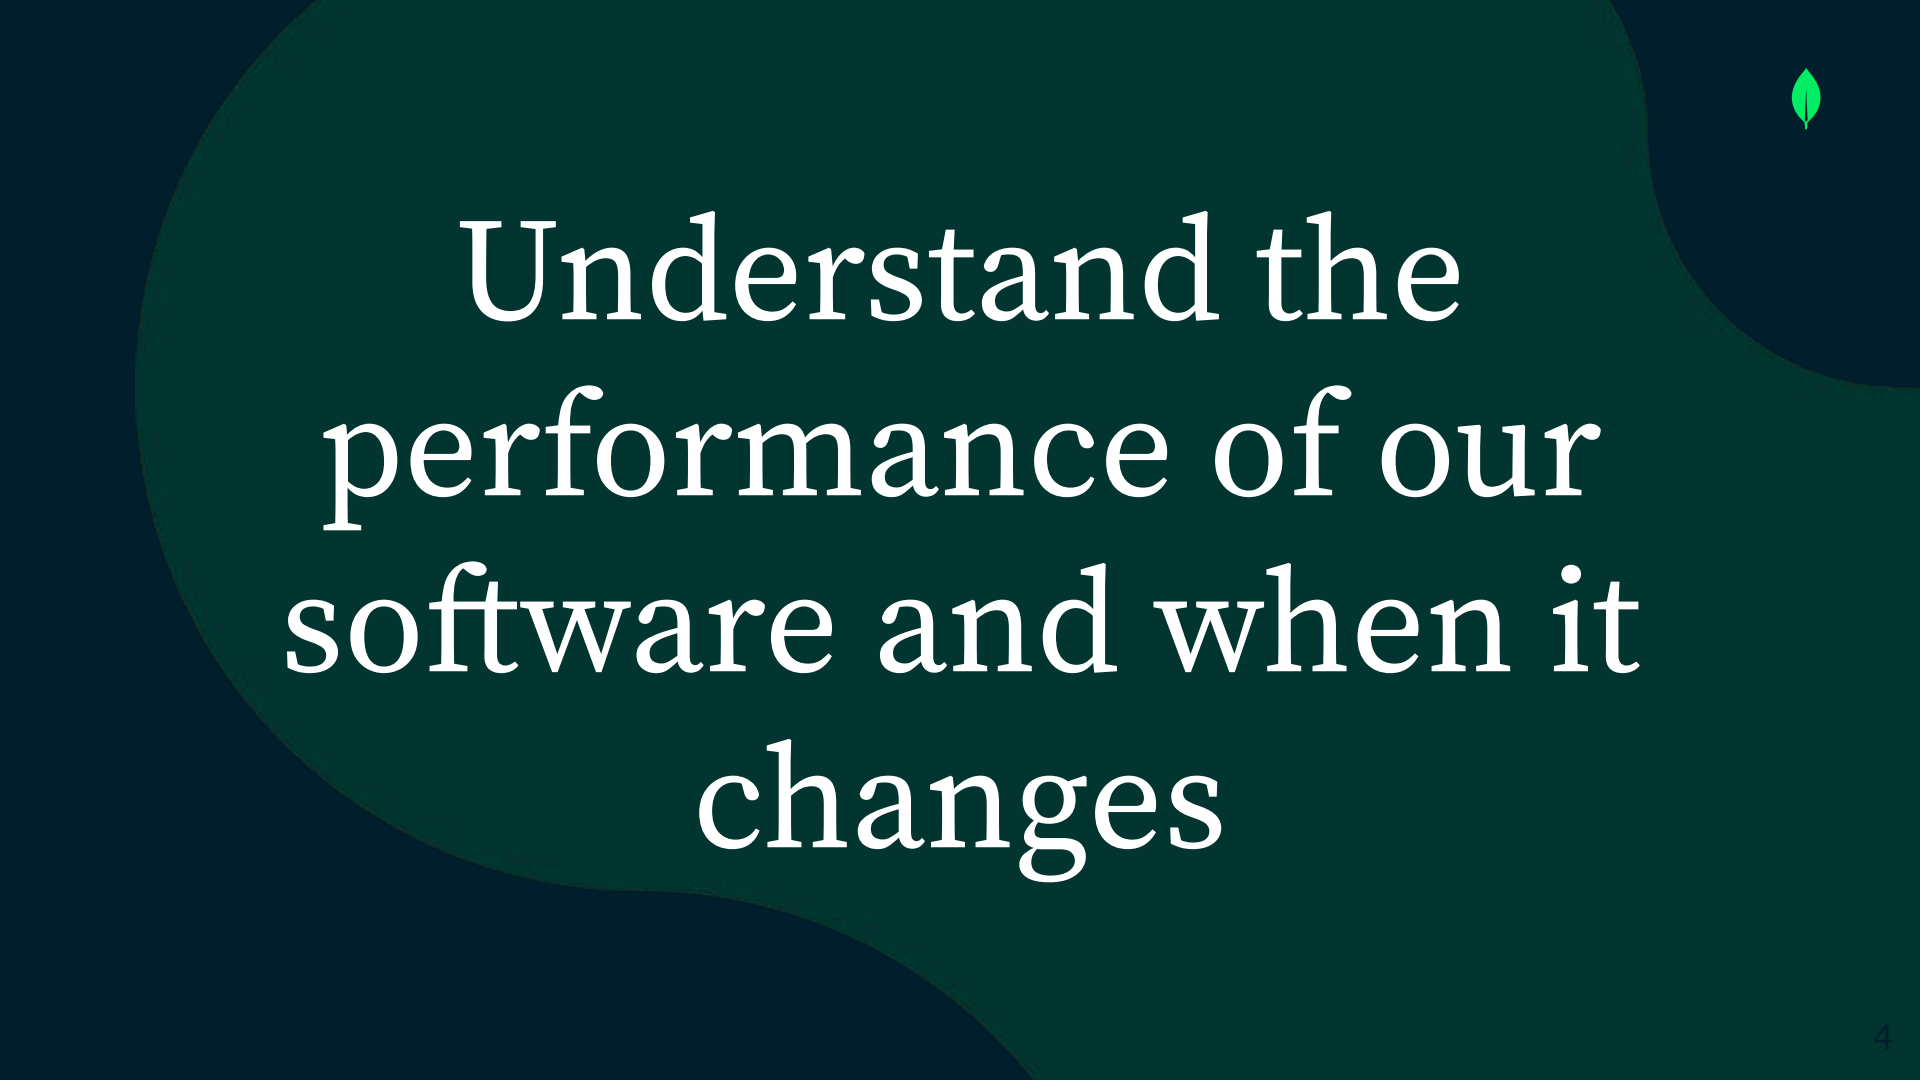
Understand the performance of our software and when it changes

The focus for DSI was serving the more complex requirements of end-to-end system performance tests on real clusters, automating every step including provisioning of hardware, and generating consistent, repeatable results

DSI Goals

- Full end-to-end automation
- Support both CI and manual testing
- Elastic, public cloud infrastructure
- Everything configurable
- All configuration via YAML
- Diagnosability
- Repeatability

DSI Modules

- Bootstrap
- Infrastructure provisioning
- System setup
- Workload setup
- MongoDB setup
- Test Control
- Analysis
- Infrastructure teardown

Configuration Files

```
mongod_config_file:
  storage:
    engine: wiredTiger
  replication:
    replSetName: rs0

topology:
  - cluster_type: replset
    id: rs0
    mongod:
      - public_ip: ${infrastructure_provisioning.out.mongod.0.public_ip}
      - public_ip: ${infrastructure_provisioning.out.mongod.1.public_ip}
      - public_ip: ${infrastructure_provisioning.out.mongod.2.public_ip}

# Meta data about this mongodb setup
meta:
  # The list of hosts that can be used in a mongodb connection string
  hosts: ${mongodb_setup.topology.0.mongod.0.private_ip}:27017
  hostname: ${mongodb_setup.topology.0.mongod.0.private_ip}
  mongodb_url: mongodb://${mongodb_setup.meta.hosts}/test?replicaSet=rs0
  is_replset: true
```

```
run:
  - id: ycsb_load
    type: ycsb
    cmd: ./bin/ycsb load mongodb -s -P ../../workloadEvergreen -threads 8
    config_filename: workloadEvergreen
    workload_config: |
      mongodb.url=${mongodb_setup.meta.mongodb_url}
      recordcount=5000000
      workload=com.yahoo.ycsb.workloads.CoreWorkload
  - id: ycsb_100read
    type: ycsb
    cmd: ./bin/ycsb run mongodb -s -P ../../workloadEvergreen_100read -threads 32
    config_filename: workloadEvergreen_100read
    workload_config: |
      mongodb.url=${mongodb_setup.meta.mongodb_url}
      recordcount=5000000
      maxexecutiontime=240
      workload=com.yahoo.ycsb.workloads.CoreWorkload
      readproportion=1.0
```


Experiments

Controlling Noise

- Repeated Test Runs (5x)
- Canary Tests + Real Tests
- Most Assumptions Were Wrong

Results

Controlling Noise

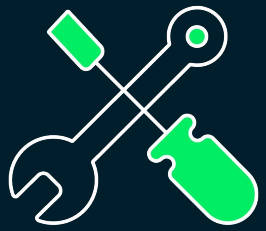

- C3.8xlarge was lowest noise
- Replaced SSD with EBS PIOPS
- Disable Optimizations

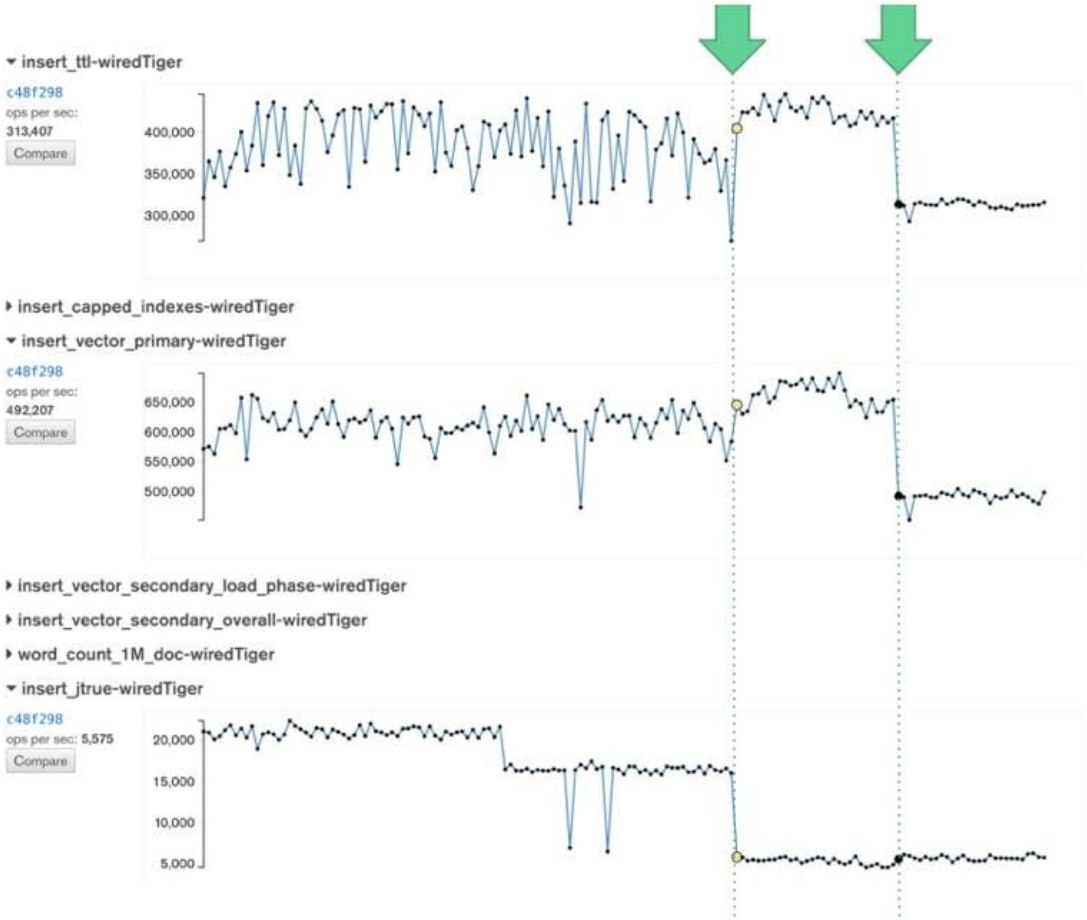

- SSD -> EBS
- CPU: No HT, single socket, scheduling

Decide if Perf Changed

Have human's look at the graphs

We have a lot of graphs

Use a threshold

High false positive and false negative rate.
Identifies regressions on the wrong commit

Use Math

Use Change Point Detection algorithms and modern math.

Problem

Problem Statement

Detect which commits change the performance of the software (as measured by our performance tests) in the presence of the noise from the testing infrastructure.

Change Point Detection

"Change point analysis is the process of detecting distributional changes within time-ordered observations."

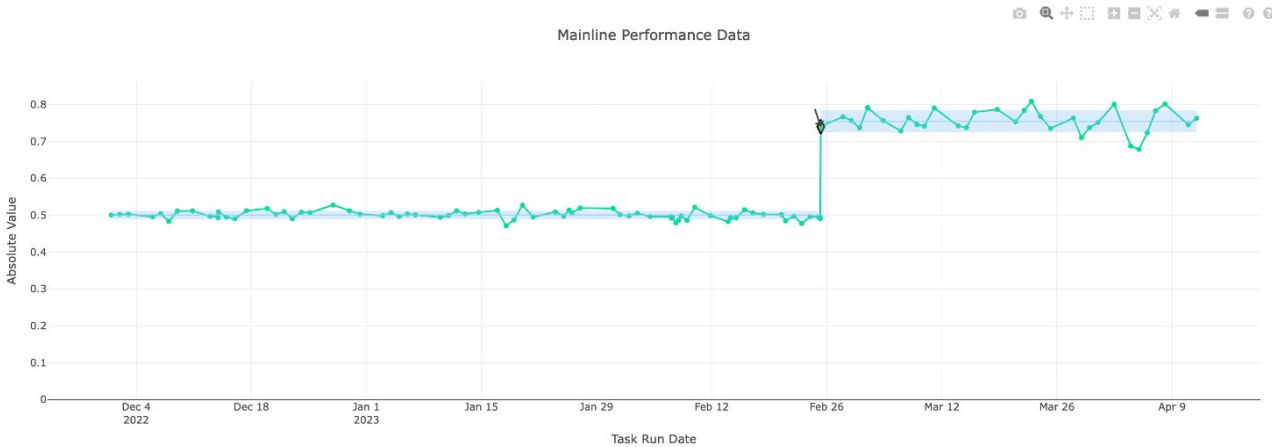

Change Point - True Positive

Project: sys-perf
Variant: linux-standalone-classic-query-engine.2022-11
Task: bestbuy_agg
Test: count_no_predicate-useAgg
Arguments: thread_level: 1
Measurement: ops_per_sec
Commit: 2f6ef2d
Commit Date: 2023-02-25T06:23:17+00:00
Previous commit: c3e2419
Previous commit date: 2023-02-25T05:12:02+00:00
JIRA Tickets: BF-27940 (BF Page)
Value: 0.7399955101396014
Absolute change: 0.25
Z Score Change: 22.96
Percentage Change: 50.79%

Triage as: NOT_TRIAGED TRUE_POSITIVE FALSE_POSITIVE UNDER_INVESTIGATION JIRA Tickets LINK TICKETS TRIAGE UNTRIAGE SCHEDULE UPDATE GRAPH SELECTION LOAD TOTAL NUMBER: FALSE Filters: COPY Paste Days Lookback 900

Graphed	Repository	Branch	Commit	Commit Dat…	Detection Date	Author	Author Email	Message	Project	Variant	Task	Test	Measurement	Arguments	Percent Change	Z-Score Change	Build Failures	Triage Status	Change Type
Not graphed	mongodb/m…	master	2f6ef2d06a…	2023-02-25…	2023-02-26…	Rui Liu	rui.liu@mon…	SERVER-74…	sys-perf	linux-microb…	agg-query-c…	Aggregation…	ops_per_sec	thread_level 1	+21.34%	+7.6	BF-27940	true_positive	improvement
Not graphed	mongodb/m…	master	2f6ef2d06a…	2023-02-25…	2023-02-26…	Rui Liu	rui.liu@mon…	SERVER-74…	sys-perf	linux-microb…	agg-query-c…	Aggregation…	ops_per_sec	thread_level 2	+20.22%	+5.35	BF-27940	true_positive	improvement
Not graphed	mongodb/m…	master	2f6ef2d06a…	2023-02-25…	2023-02-26…	Rui Liu	rui.liu@mon…	SERVER-74…	sys-perf	linux-microb…	agg-query-c…	Aggregation…	ops_per_sec	thread_level 8	+19.8%	+6.85	BF-27940	true_positive	improvement
Not graphed	mongodb/m…	master	2f6ef2d06a…	2023-02-25…	2023-02-26…	Rui Liu	rui.liu@mon…	SERVER-74…	sys-perf	linux-microb…	agg-query-c…	Aggregation…	ops_per_sec	thread_level 4	+22.55%	+4.72	BF-27940	true_positive	improvement
Not graphed	mongodb/m…	master	2f6ef2d06a…	2023-02-25…	2023-02-28…	Rui Liu	rui.liu@mon…	SERVER-74…	sys-perf	linux-microb…	agg-query-c…	Aggregation…	ops_per_sec	thread_level 2	+21.16%	+4.69	BF-27940	true_positive	improvement
Not graphed	mongodb/m…	master	2f6ef2d06a…	2023-02-25…	2023-02-26…	Rui Liu	rui.liu@mon…	SERVER-74…	sys-perf	linux-microb…	agg-query-c…	Aggregation…	ops_per_sec	thread_level 1	+17.99%	+5.17	BF-27940	true_positive	improvement
Not graphed	mongodb/m…	master	2f6ef2d06a…	2023-02-25…	2023-02-26…	Rui Liu	rui.liu@mon…	SERVER-74…	sys-perf	linux-microb…	agg-query-c…	Aggregation…	ops_per_sec	thread_level 8	+17.93%	+5.06	BF-27940	true_positive	improvement
#00de9b	mongodb/m…	master	2f6ef2d06a…	2023-02-25…	2023-02-26…	Rui Liu	rui.liu@mon…	SERVER-74…	sys-perf	linux-standal…	bestbuy_agg	count_no_pr…	ops_per_sec	thread_level 1	+50.79%	+22.96	BF-27940	true_positive	improvement
Not graphed	mongodb/m…	master	2f6ef2d06a…	2023-02-25…	2023-02-26…	Rui Liu	rui.liu@mon…	SERVER-74…	sys-perf	linux-microb…	views-aggre…	Aggregation…	ops_per_sec	thread_level 1	+75.65%	+26.55	BF-27940	true_positive	improvement
Not graphed	mongodb/m…	master	2f6ef2d06a…	2023-02-25…	2023-02-26…	Rui Liu	rui.liu@mon…	SERVER-74…	sys-perf	linux-microb…	views-aggre…	Aggregation…	ops_per_sec	thread_level 1	+17.17%	+5.16	BF-27940	true_positive	improvement

Page 1 of 20 Results per page 50 Showing 1 - 50 of 1,000

Build Failures / BF-27940

52% improvement in ops_per_sec in bestbuy_agg/ count_no_predicate-useAgg (System Performance) - 2f6ef2d, 2023-02-25

Edit | Add comment | Assign | More | Reopen Issue | Flag for Baron Triage

Details

Type:	Build Failure
Priority:	Major - P3
Labels:	bf-code perf-improvement
Sprint:	QE 2023-03-20
Severity Type:	Not Release Blocking
De-escalation Justification:	De-escalated by the Build Baron team
Evergreen Project:	sys-perf
Failing Buildvariants:	linux-standalone-classic-query-engine.2022-11
Failing Tasks:	bestbuy_agg
Failing Tests:	count_no_predicate-useAgg
First Failing Revision:	2f6ef2d06acc9b925c7c34ca66bcc16a1fe85ead
Score:	35
Score Breakdown:	FREQUENCY SCORE: 0. Fewer than 2 recent Evergreen executions. RECENCY SCORE: 35. First known occurrence was today.
Linked BFG List:	https://performance-monitoring-and-analysis.server-tig.prod.corp.mongodb.com/bf/BF-27940
Type of fix:	No code change
Performance Change Type:	Improvement
Status:	CLOSED
Resolution:	Works as Designed

People

Assignee:	Rui Liu Assign to me
Reporter:	Sviatlana Zuiko
Votes:	0 Vote for this issue
Watchers:	2 Start watching this issue

Dates

Created:	Mar 03 2023 08:29:36 PM EST
Updated:	Mar 05 2023 01:03:54 PM EST
Resolved:	Mar 04 2023 02:05:19 PM EST
Team Assignment:	04/Mar/23 2:04 PM

Agile

Completed Sprint: QE 2023-03-20 ended 17/Mar/23
View on Board

Reporter

name:	sviatlana.zuiko@mongodb.com
email:	sviatlana.zuiko@mongodb.com

Description

52% improvement in ops_per_sec in bestbuy_agg/ count_no_predicate-useAgg

THIS JIRA TICKET HAS PRELIMINARY INFO ONLY and is considered to be kind of a placeholder for a Perf BF. Please visit **Perf BF** page, which contains a consolidated performance impact, for more details.

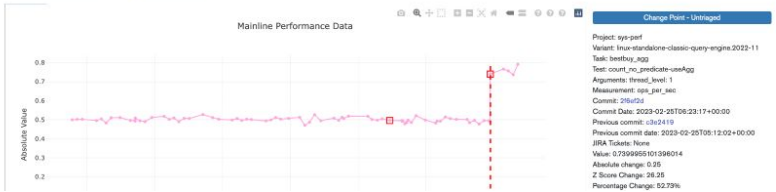

BF-27940

CHANGE POINTS | FAILING PROJECTS | FAILING VARIANTS | FAILING TASKS | FAILING TESTS | FAILING MEASUREMENTS | FAILING ARGUMENTS

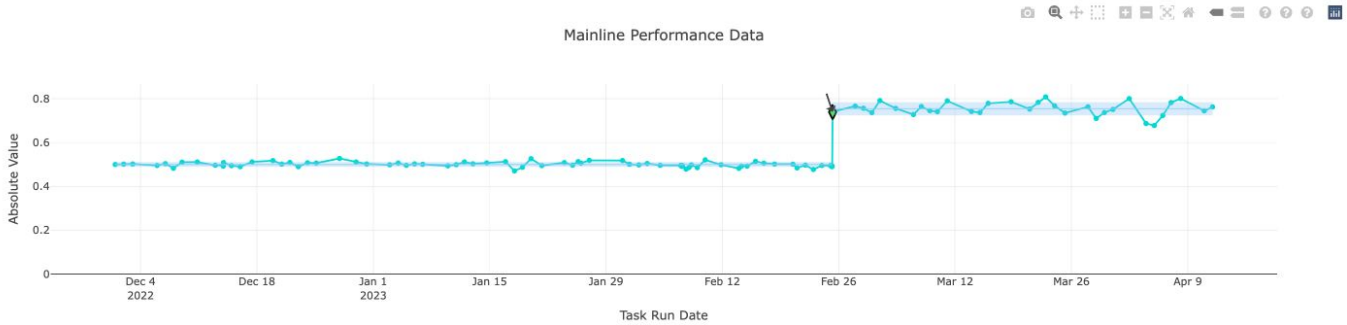

Change Points:

GRAPH MAINLINE PERF DATA

Graphed	Repository…	Branch	Commit	Commit Dat…	Detection D…	Author	Author Emai…	Message	Project	Variant	Task	Test	Measureme…	Arguments…	Percent Ch…	Z-Score Ch…	Build Failur…	Triage Statu…	Change Typ…
Not graphed	mongodb/m…	master	2f6ef2d06a…	2023-02-2…	2023-02-2…	Rui Liu	rui.liu@mon…	SERVER-7…	sys-perf	linux-microb…	views-aggre…	Aggregation…	ops_per_sec	thread_level 1	75.646811…	26.547426…	BF-27940	true_positive	improvement
#00dad2	mongodb/m…	master	2f6ef2d06a…	2023-02-2…	2023-02-2…	Rui Liu	rui.liu@mon…	SERVER-7…	sys-perf	linux-standa…	bestbuy_agg	count_no_p…	ops_per_sec	thread_level 1	50.790886…	22.959368…	BF-27940	true_positive	improvement
Not graphed	mongodb/m…	master	2f6ef2d06a…	2023-02-2…	2023-02-2…	Rui Liu	rui.liu@mon…	SERVER-7…	sys-perf	linux-microb…	agg-query-c…	Aggregation…	ops_per_sec	thread_level 4	22.552994…	4.7197364…	BF-27940	true_positive	improvement
Not graphed	mongodb/m…	master	2f6ef2d06a…	2023-02-2…	2023-02-2…	Rui Liu	rui.liu@mon…	SERVER-7…	sys-perf	linux-microb…	agg-query-c…	Aggregation…	ops_per_sec	thread_level 4	21.389287…	7.0951588…	BF-27940	true_positive	improvement
Not graphed	mongodb/m…	master	2f6ef2d06a…	2023-02-2…	2023-02-2…	Rui Liu	rui.liu@mon…	SERVER-7…	sys-perf	linux-microb…	agg-query-c…	Aggregation…	ops_per_sec	thread_level 1	21.336124…	7.6033680…	BF-27940	true_positive	improvement
Not graphed	mongodb/m…	master	2f6ef2d06a…	2023-02-2…	2023-02-2…	Rui Liu	rui.liu@mon…	SERVER-7…	sys-perf	linux-microb…	agg-query-c…	Aggregation…	ops_per_sec	thread_level 2	21.162847…	4.6915122…	BF-27940	true_positive	improvement
Not graphed	mongodb/m…	master	2f6ef2d06a…	2023-02-2…	2023-02-2…	Rui Liu	rui.liu@mon…	SERVER-7…	sys-perf	linux-microb…	agg-query-c…	Aggregation…	ops_per_sec	thread_level 2	20.224006…	5.3529406…	BF-27940	true_positive	improvement

Scaling Our Testing

New Tools

- Load Generator – Measure Everything
- More Metrics
- Measure System Metrics

[53]:
```
plot(li512)
```

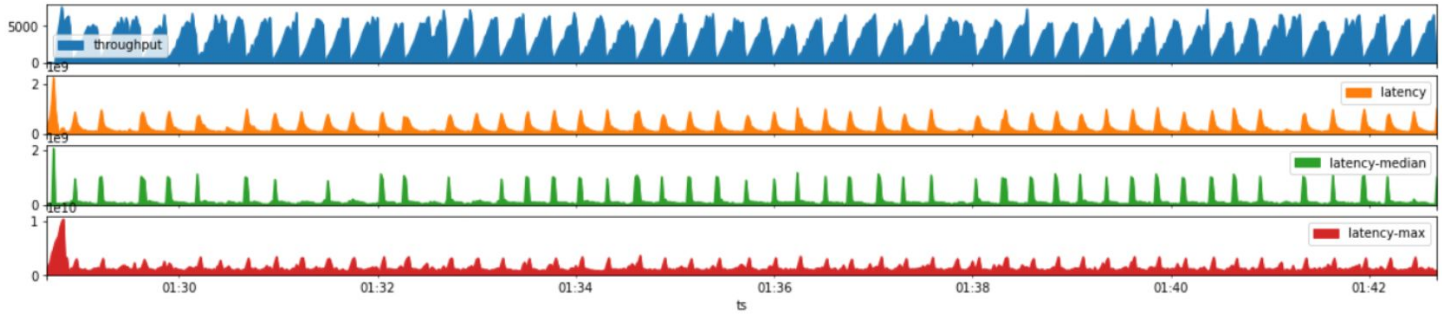

[54]:
```
plot(su512)
```

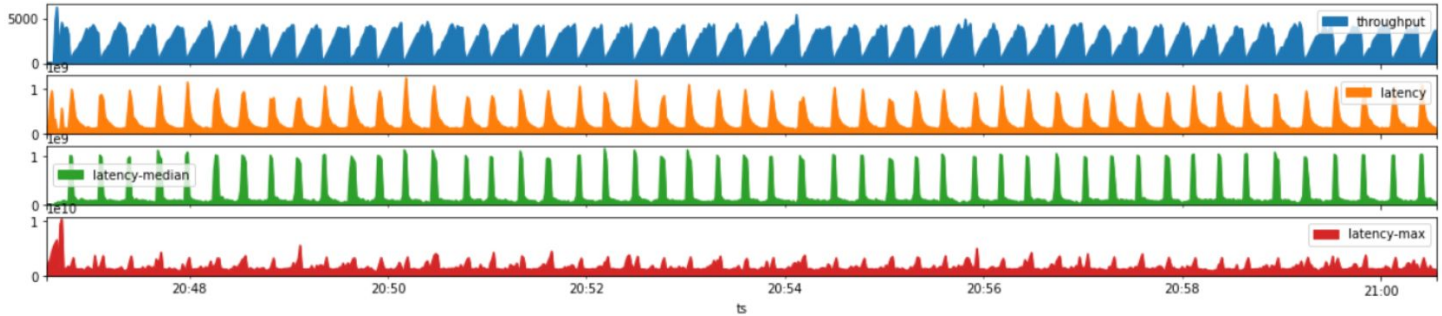

Problem is getting harder

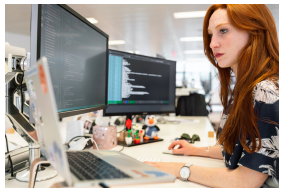

x

x

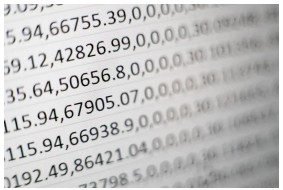

=

It Keeps Growing 5-10x a year!

	Variants	Tasks	Results
2018	10	80	
2019	17	170	2,195
2020	24	342	17,240
2021	30	395	157,834
2022	39	655	599,130
recent	54	1,054	1,051,150

Add Process

Throw People At the Problem

- Dedicated Build Barons
- Monthly Review of Performance
- Try to Automate More

Adding Features + Updates

New Dashboards

Compare experiment to mainline

Compare a branch/or feature on to disabled (SBE build)

All enabled from existing, but requiring more work

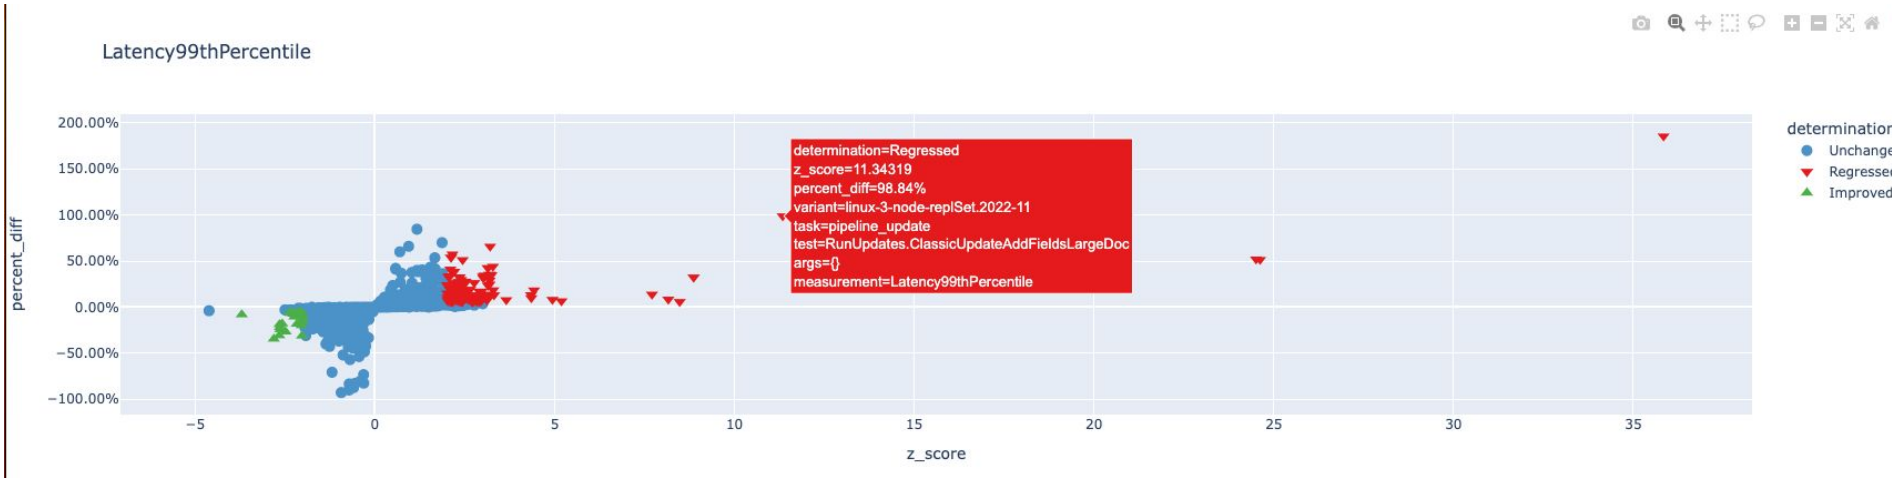
Latency99thPercentile
200.00%
150.00%
100.00%
50.00%
0.00%
−50.00%
−100.00%
percent_diff
−5
0
5
10
15
20
25
30
35
z_score
determination=Regressed
z_score=11.34319
percent_diff=98.84%
variant=linux-3-node-replSet.2022-11
task=pipeline_update
test=RunUpdates.ClassicUpdateAddFieldsLargeDoc
args={}
measurement=Latency99thPercentile
determination
Unchange
Regressed
Improved


63d2a2733627e01fb8446106 : linux-3-node-replSet.2022-11 : industry_benchmarks : ycsb_50read50update : average_read_latency_us : {"thread
average_read_latency_us
320
300
280
260
240
220
Feb 19
2023
Feb 26
Mar 5
Mar 12
Mar 19
Mar 26
Apr 2
Apr 9
Apr 16
Commit Date
Legend
Base Timeseries
Your Value
Base Value
Base Date
Stable Average
Stable Region
Stable Region Length: 9
Z-Score: -3.6
Percent Diff (Region): -10.01%


PM-2697: SBE Perf / Milestone Progress
Resolved issues ratio
93%
43
46
M1 (PERF INFRA)
49%
28
57
M2 (CORE ARCH)
36%
19
53
M3 (TARGET WORKLOADS)
52%
22
42
M4 (OPTIONALS)
207 total issues
Issue Count by Milestones


PM-2697: SBE Perf / Milestone Progress
Resolved issues ratio
100%
46
46
M1 (PERF INFRA)
96%
52
54
M2 (CORE ARCH)
61%
37
61
M3 (TARGET WORKLOADS)
39%
29
75
M4 (OPTIONALS)
262 total issues
Issue Count by Milestones

Feature Dashboard

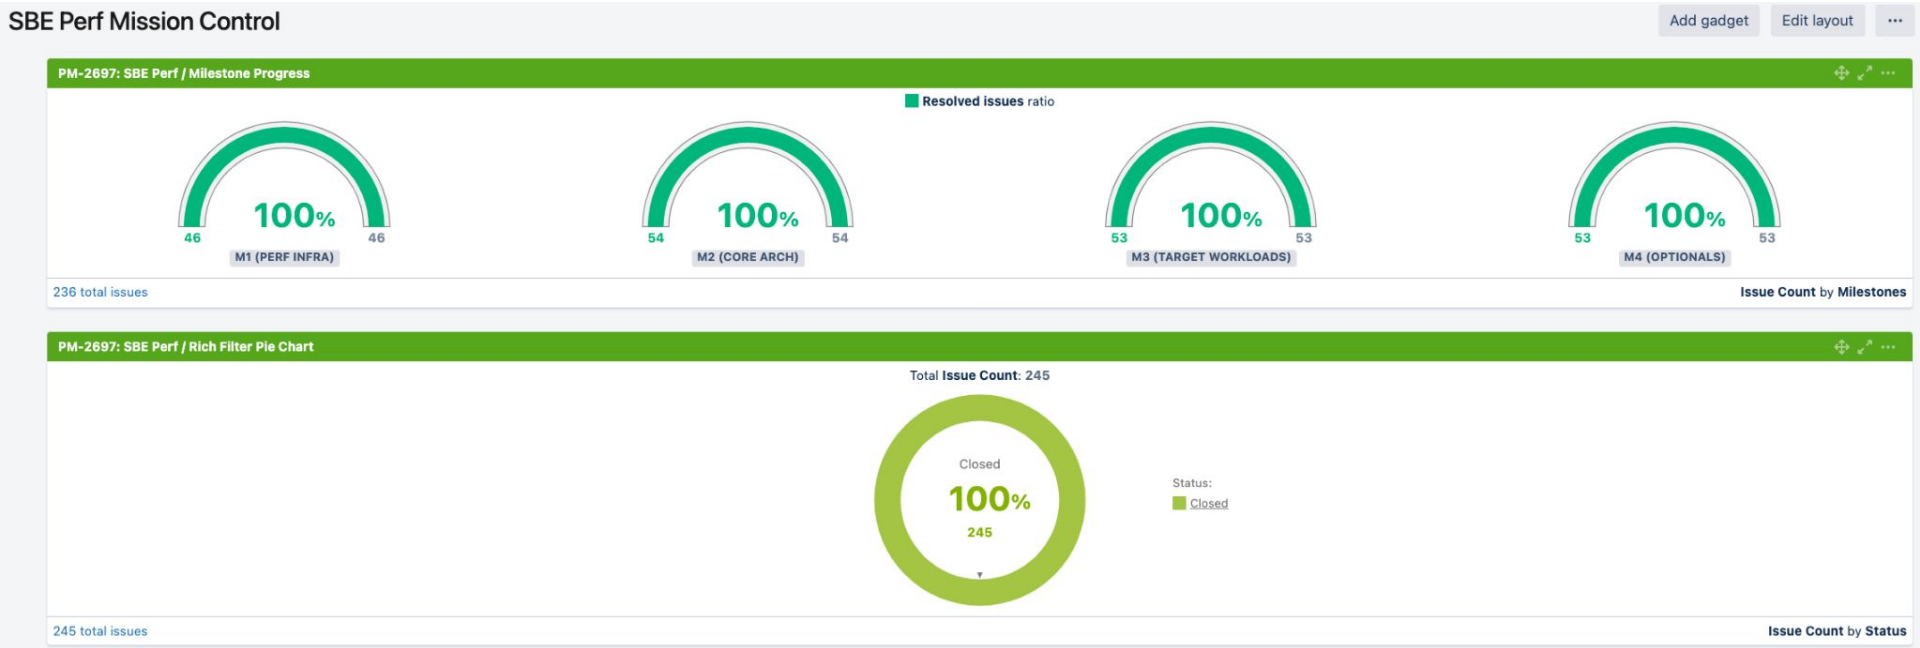

Updating Test Systems

Pinned Systems

Systems became old!

Must Control Noise

Ran several extensive experiments

How to Upgrade

Added versioning to our system. Add new configurations – phase out old ones

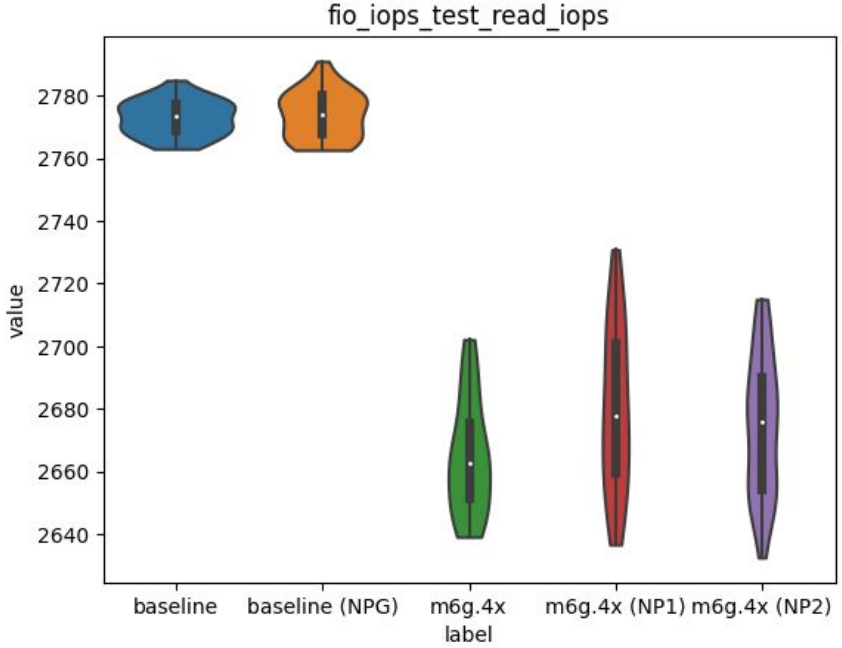
fio_iops_test_read_iops
value
2780
2760
2740
2720
2700
2680
2660
2640
baseline
baseline (NPG)
m6g.4x
m6g.4x (NP1)
m6g.4x (NP2)
label

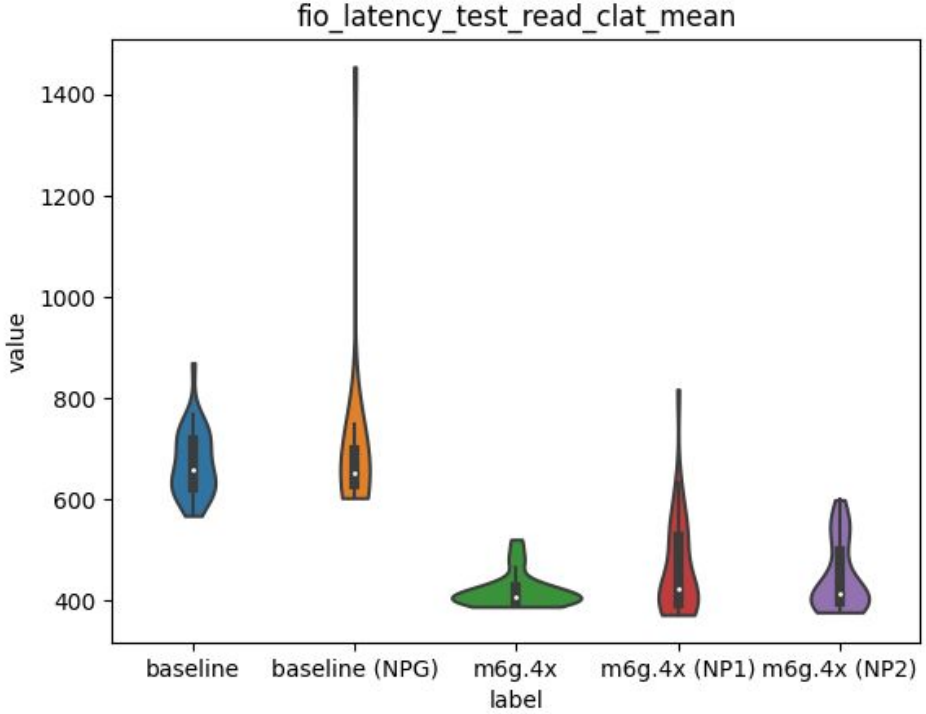
fio_latency_test_read_clat_mean
value
1400
1200
1000
800
600
400
baseline
baseline (NPG)
m6g.4x
m6g.4x (NP1)
m6g.4x (NP2)
label

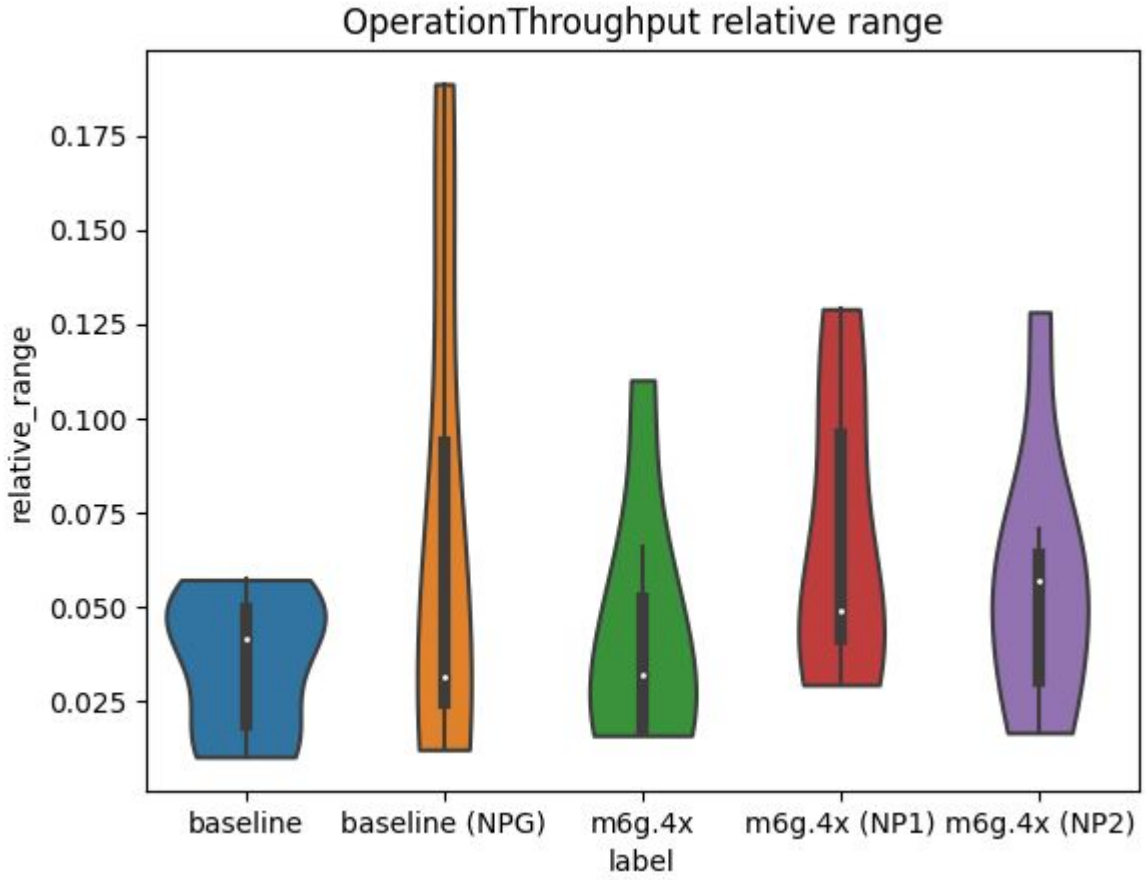
OperationThroughput relative range
relative_range
0.175
0.150
0.125
0.100
0.075
0.050
0.025
baseline
baseline (NPG)
m6g.4x
m6g.4x (NP1)
m6g.4x (NP2)
label

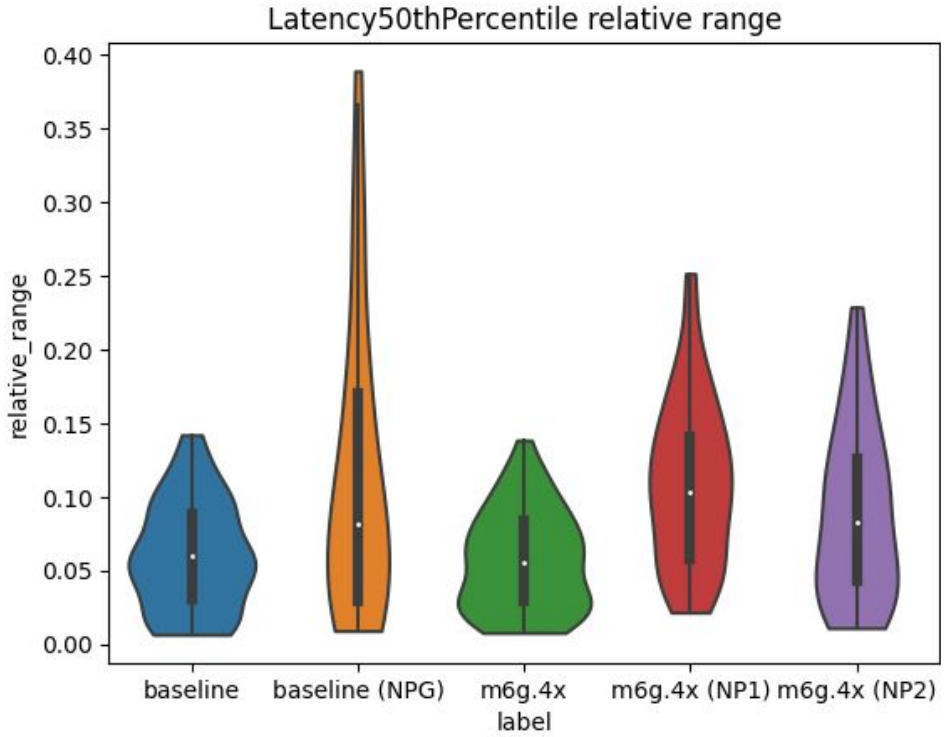
Latency50thPercentile relative range
relative_range
0.40
0.35
0.30
0.25
0.20
0.15
0.10
0.05
0.00
baseline
baseline (NPG)
m6g.4x
m6g.4x (NP1)
m6g.4x (NP2)
label

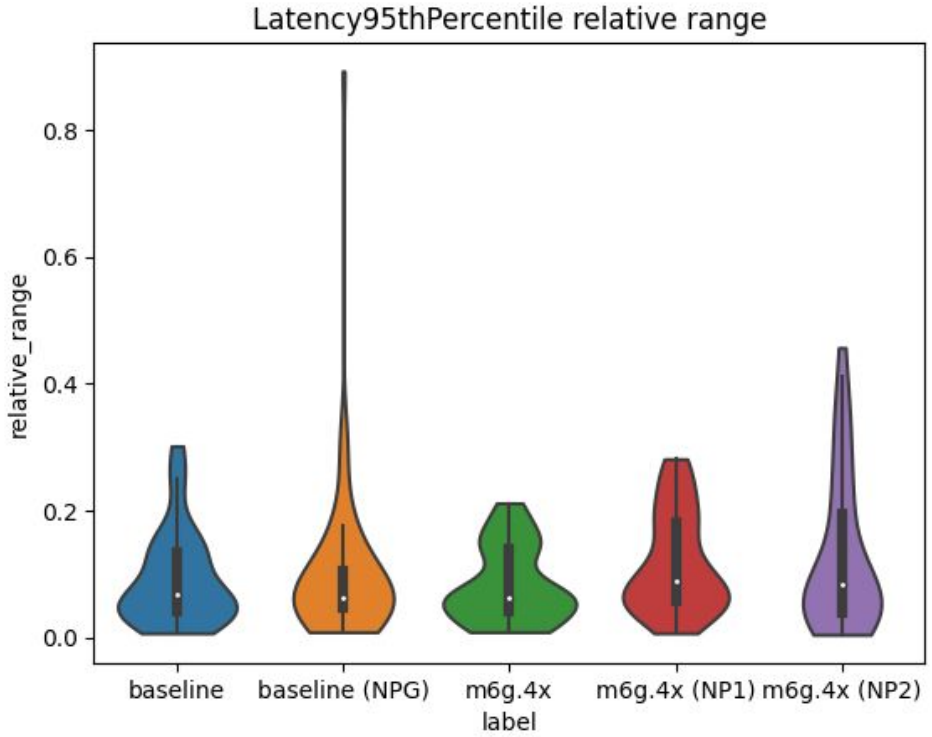
Latency95thPercentile relative range
relative_range
0.8
0.6
0.4
0.2
0.0
baseline
baseline (NPG)
m6g.4x
m6g.4x (NP1)
m6g.4x (NP2)
label

Diagnosis

Visualization of telemetry

Profiling support

FTDC(t2)

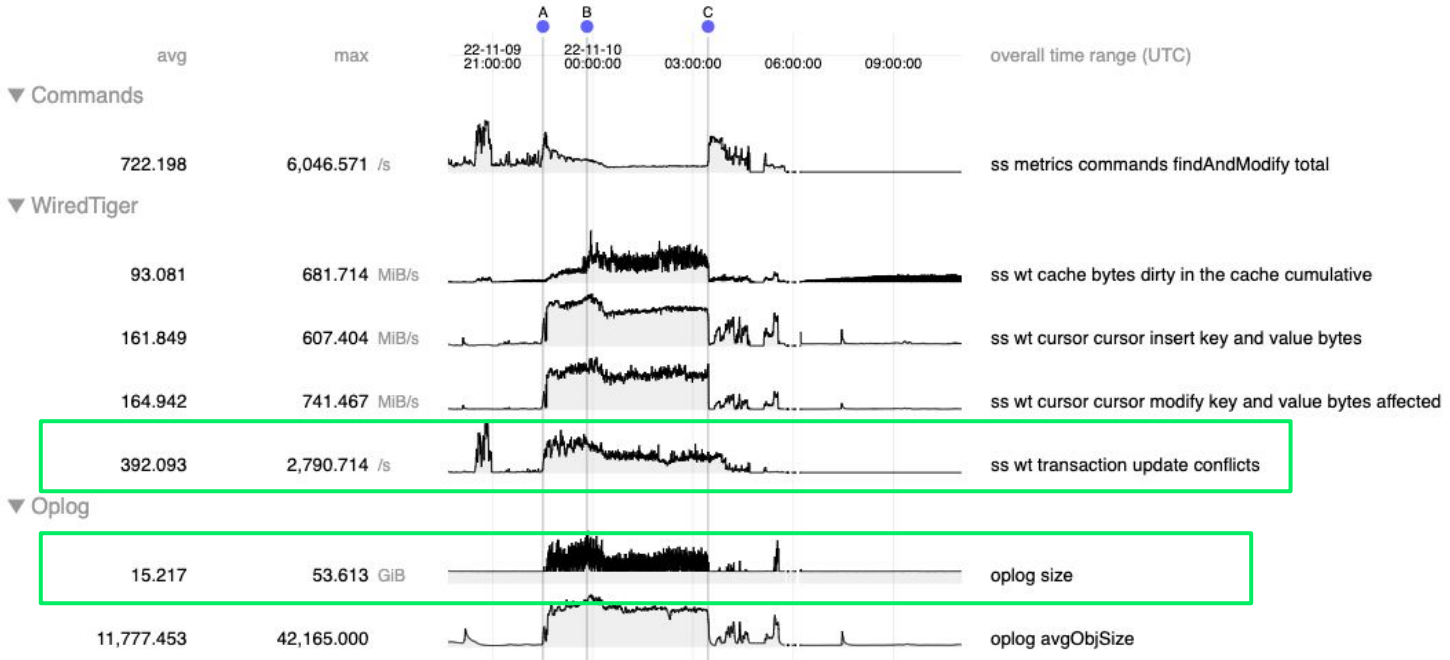
A
B
C
avg
max
22-11-09 21:00:00
22-11-10 00:00:00
03:00:00
06:00:00
09:00:00
overall time range (UTC)
▼ Commands
722.198
6,046.571 /s
ss metrics commands findAndModify total
▼ WiredTiger
93.081
681.714 MiB/s
ss wt cache bytes dirty in the cache cumulative
161.849
607.404 MiB/s
ss wt cursor cursor insert key and value bytes
164.942
741.467 MiB/s
ss wt cursor cursor modify key and value bytes affected
392.093
2,790.714 /s
ss wt transaction update conflicts
▼ Oplog
15.217
53.613 GiB
oplog size
11,777.453
42,165.000
oplog avgObjSize

Profiling with Flamegraphs

- Can automatically rerun task to
  - Profile
  - Create Flame Graph

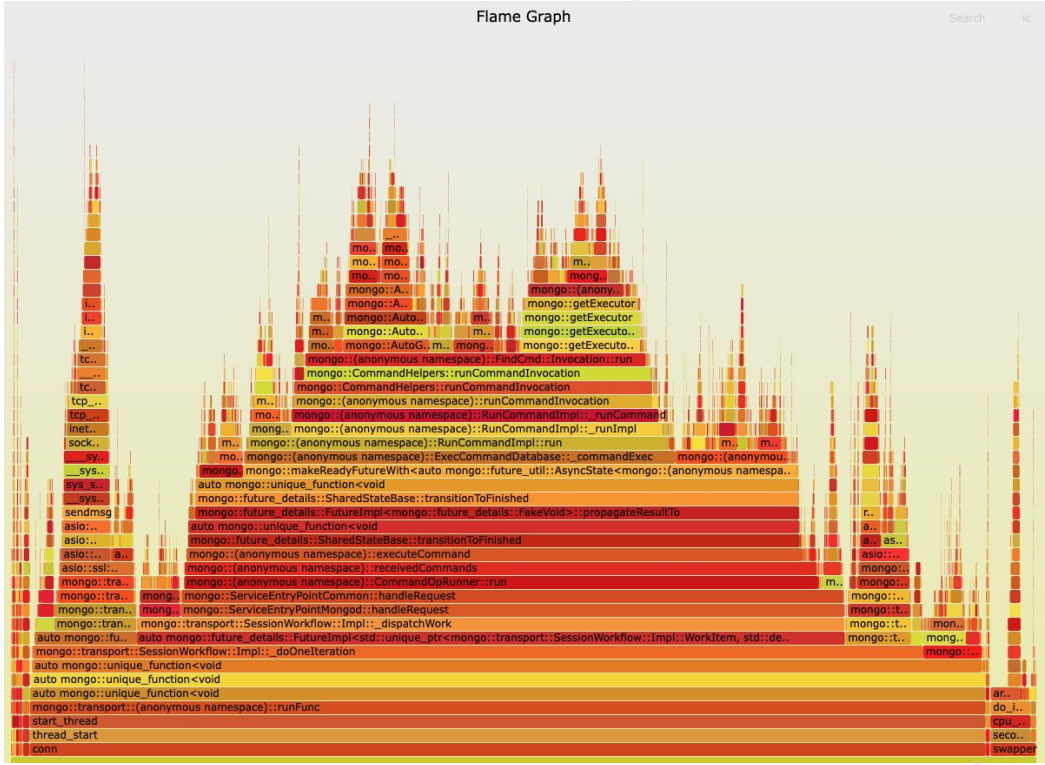

[Brendan Gregg FlameGraphs]

Stability >> Raw Performance

How do you measure stability?

Not well

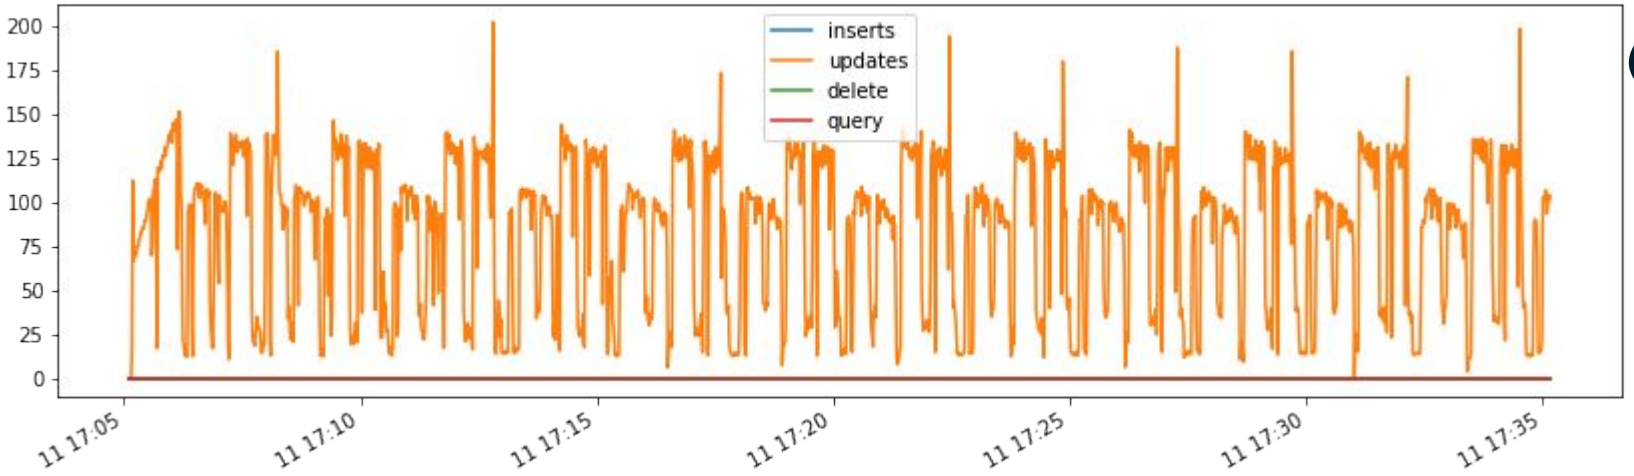
inserts
updates
delete
query
200
175
150
125
100
75
50
25
0
11 17:05
11 17:10
11 17:15
11 17:20
11 17:25
11 17:30
11 17:35

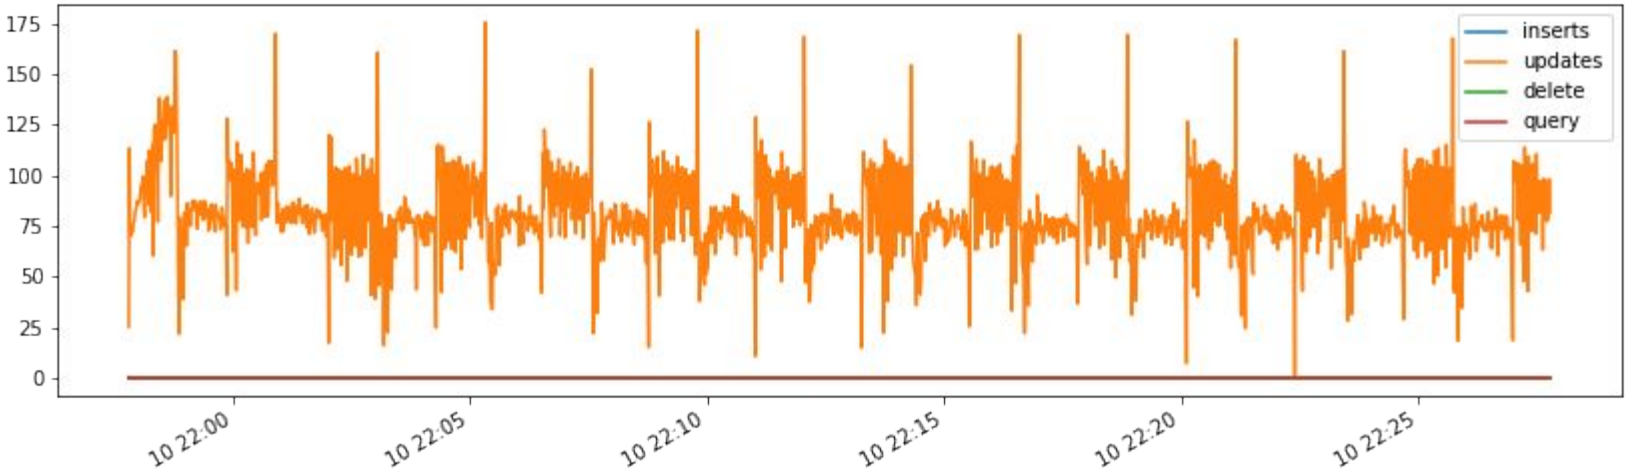
inserts
updates
delete
query
175
150
125
100
75
50
25
0
10 22:00
10 22:05
10 22:10
10 22:15
10 22:20
10 22:25

Community Interactions

Data Challenge

- CFP
- Notes
- Data
- 4 Accepted Papers

SPEC-RG

Hunter Paper Slide

Hunter: Using Change Point Detection to Hunt for Performance Regressions

Matt Fleming*
matt@codeblueprint.co.uk

Piotr Kołaczkowski*
pkolaczk@datastax.com

Ishita Kumar*
ishitakumar@umass.edu

Shaunak Das*
shaunak.das@datastax.com

Sean McCarthy*
sean.mccarthy@datastax.com

Pushkala Pattabhiraman*
Pushkala.Pattabhiraman@datastax.com

Henrik Ingo*
henrik.ingo@avoinelama.fi

ABSTRACT

Change point detection has recently gained popularity as a method of detecting performance changes in software due to its ability to cope with noisy data. In this paper we present Hunter, an open source tool that automatically detects performance regressions and improvements in time-series data. Hunter uses a modified E-divisive means algorithm to identify statistically significant changes in normally-distributed performance metrics. We describe the changes we made to the E-divisive means algorithm along with their motiva-

of ICPE 2023: International Conference on Performance Engineering (ICPE2023). ACM, New York, NY, USA, 8 pages. https://doi.org/10.1145/1122445.1122456

1 INTRODUCTION

Testing the performance of distributed databases, such as Apache Cassandra, is an integral part of the development process and is often incorporated into Continuous Integration pipelines where performance tests and benchmarks can be run periodically or in response to pushing changes to source code repositories. But given

Work With Us

We have real world problems and would love to work with the community

- Noise Reduction work
- Dbtest.io: "Automated System Performance Testing at MongoDB"
- ICPE 2020: "The Use of Change Point Detection to Identify Software Performance Regressions in a Continuous Integration System" (video)
- ICPE 2021: Creating a Virtuous Cycle in Performance Testing at MongoDB

Our code is open source: signal-processing-algorithms, infrastructure code

Final Thoughts

For academics

Consider the practical implications
Humans have to use these things and they have to scale
There are always things we need to work on. We don't need to find the optimal solution for anything, just a good enough solution for the most important things.

Practitioners

Automate everything. Consider the humans. Leverage what's out there, and share what you learn

Thank You