

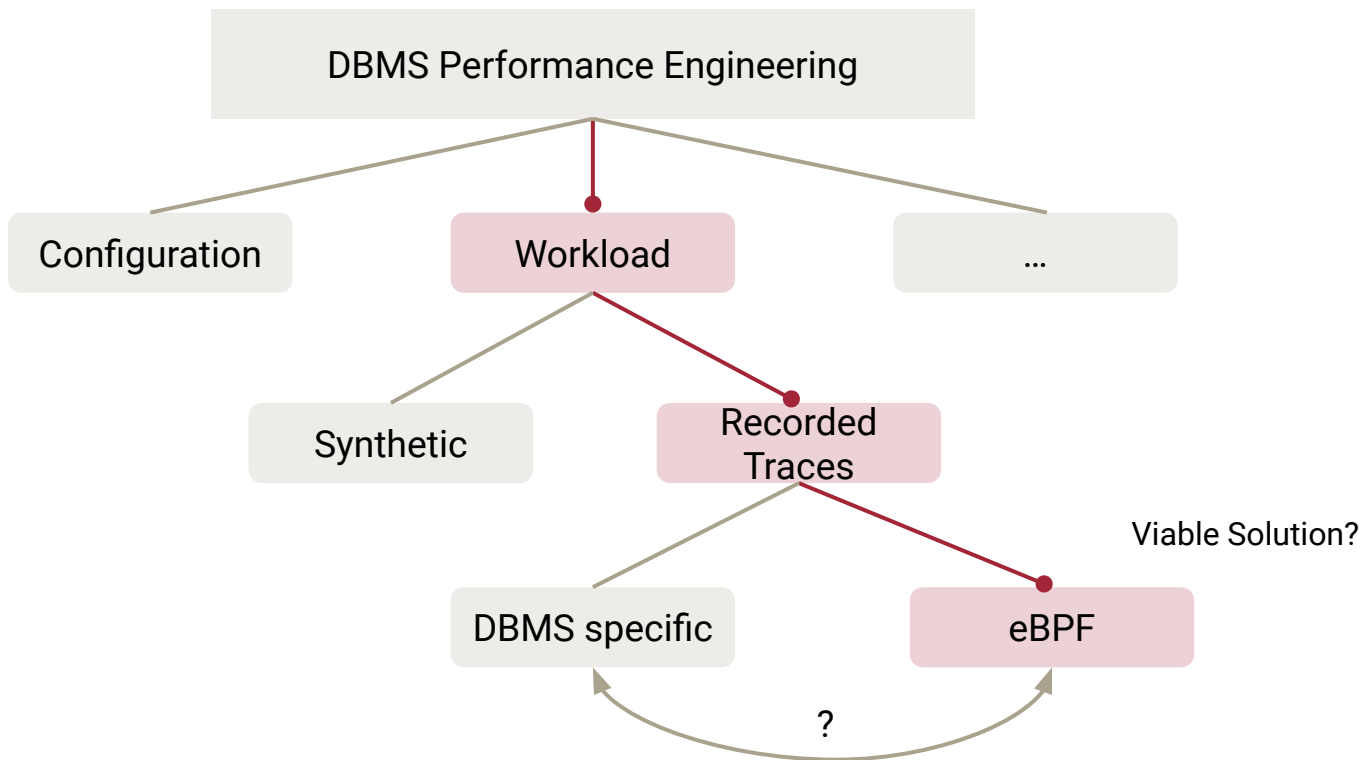


# Using eBPF for Database Workload Tracing: An Explorative Study



Georg Eisenhart  
Jörg Domaschka, Simon Volpert, Kevin Maier, Georg Eisenhart, Daniel Seybold

# Motivation

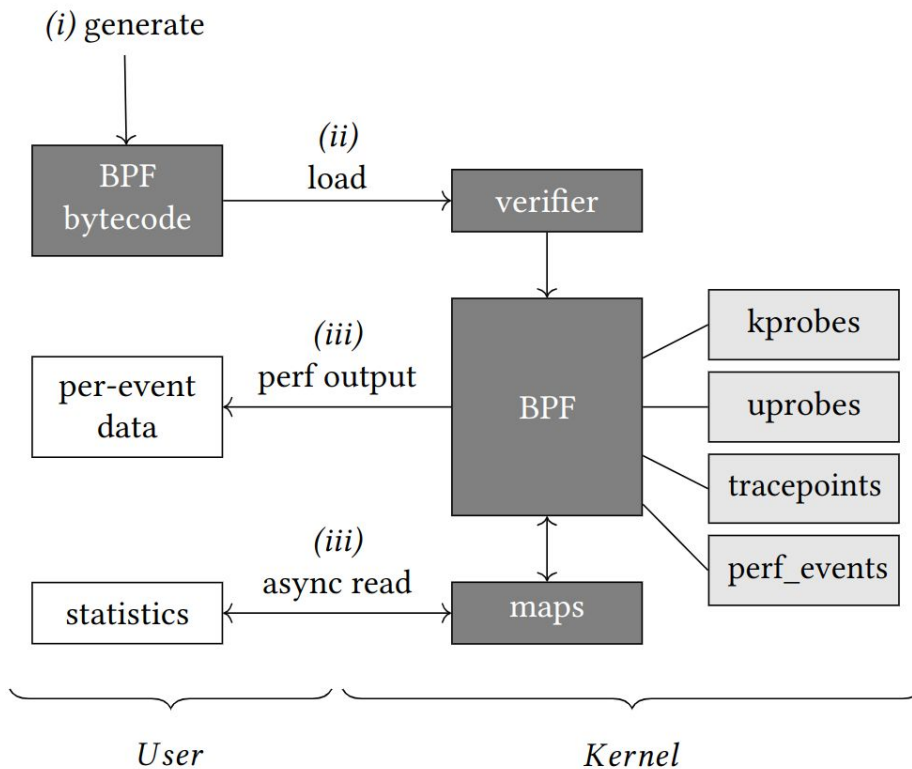


## Research Questions

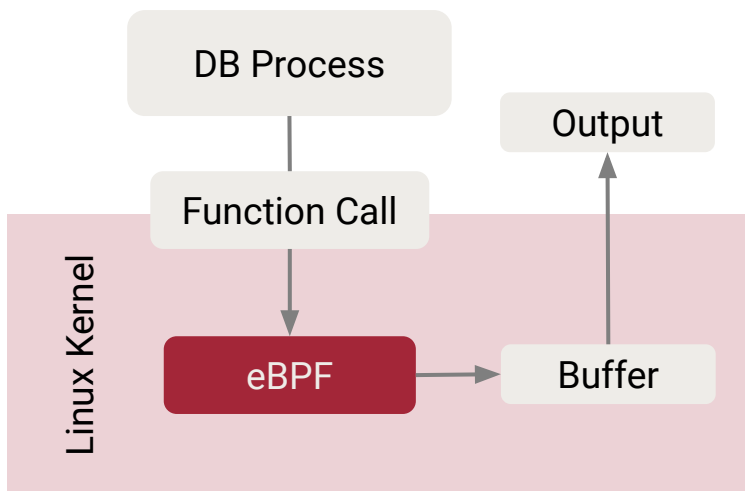
eBPF aims to enable low-overhead observability & tracing

- ▶ **RQ1** Can DBMS be instrumented (using eBPF) in order to trace occurring workload?
- ▶ **RQ2** How big is the impact of such an (eBPF) instrumentation on the overall performance?
- ▶ **RQ3** How does the eBPF impact compare to native DBMS tracing?

## eBPF Overview



## eBPF Example



```
bpf.attach_uprobe(name=args.path,
sym=="\\w+dispatch_command\\w+", fn_name="query_start")
```

```
int query_start(struct pt_regs *ctx) {
    int zero = 0;
    struct temp_t *tmp = temp_data_buffer.lookup(&zero);
    if (!tmp)
        return 0;
    tmp->timestamp = bpf_ktime_get_ns();

    #if defined(MYSQL56)
        bpf_probe_read_user(&tmp->query, sizeof(tmp->query),
        (void*) PT_REGS_PARM3(ctx));
        . . .
    #endif

    u64 pid = bpf_get_current_pid_tgid();
    temp_data_buffer.update(&zero, tmp);
    return 0;
}
```

## Probe Selection

- ▶ User space tracing
  - ▶ USDT and uprobes allow capturing function invocations including parameters and execution time.
- ▶ We focus on **uprobes**
- ▶ Attach uprobe to the query start
- ▶ Attach uretprobe to query end

	static	dynamic	userspace	kernel-space
counter	✓	✗	✗	✓
tracepoint	✓	✗	✗	✓
kprobe	✗	✓	✗	✓
uprobe	✗	✓	✓	✗
USDT	✓	✗	✓	✗

## Toolset

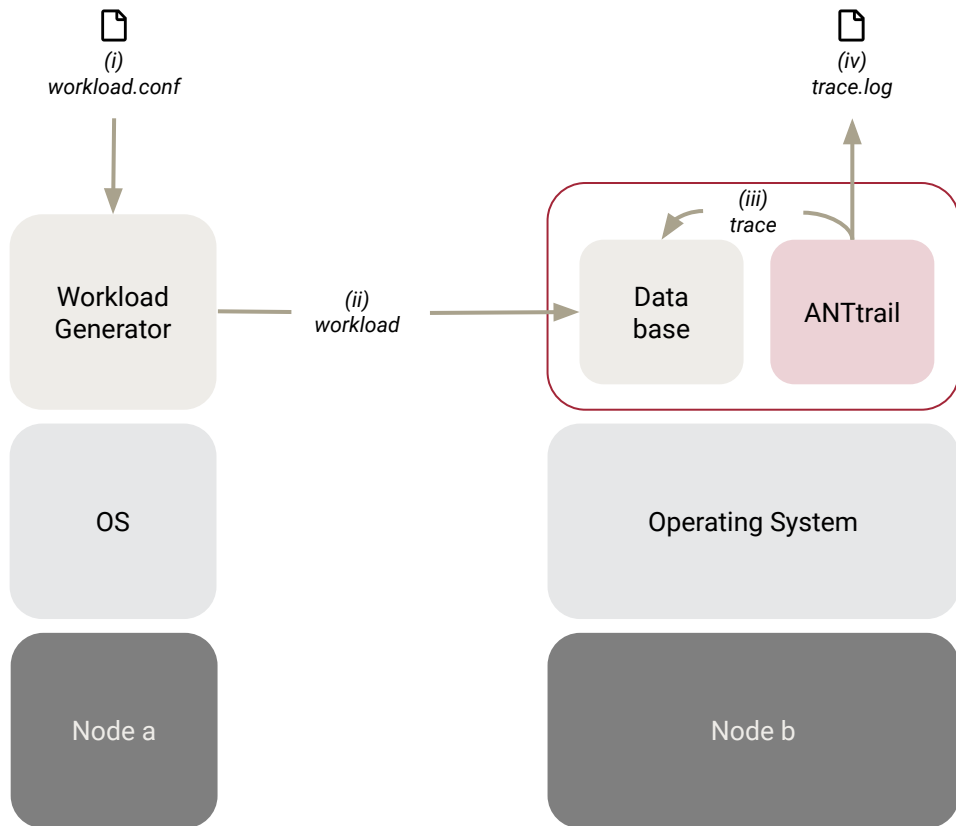
- ▶ **Baseline tool**
  - ▶ Dbslower<sup>1</sup>
  - ▶ On function call detection collect timestamp and arguments
  - ▶ On end detection calculate execution time
- ▶ **Tool Extensions**
  - ▶ Extended char array to hold full queries
  - ▶ Added support for PostgreSQL uprobes
  - ▶ Tracing and logging of any query can be written to files
  - ▶ Ability to replay later for production workload analysis
  - ▶ Remove filters to be able to trace fast queries



**ANTtrail**

1) <https://github.com/iovisor/bcc/blob/master/tools/dbslower.py>

## Experiment Workflow





# Evaluation Scenarios

## 1. Baseline

DBMS Performance without any tracing in place

## 2. DBMS Native

Native DBMS specific tracing capabilities enabled

## 3. eBPF Active

Detect queries with ANTtrail, without processing

## 4. eBPF Process

+

Process queries with ANTtrail, but not persisting

## 5. eBPF Persist

+

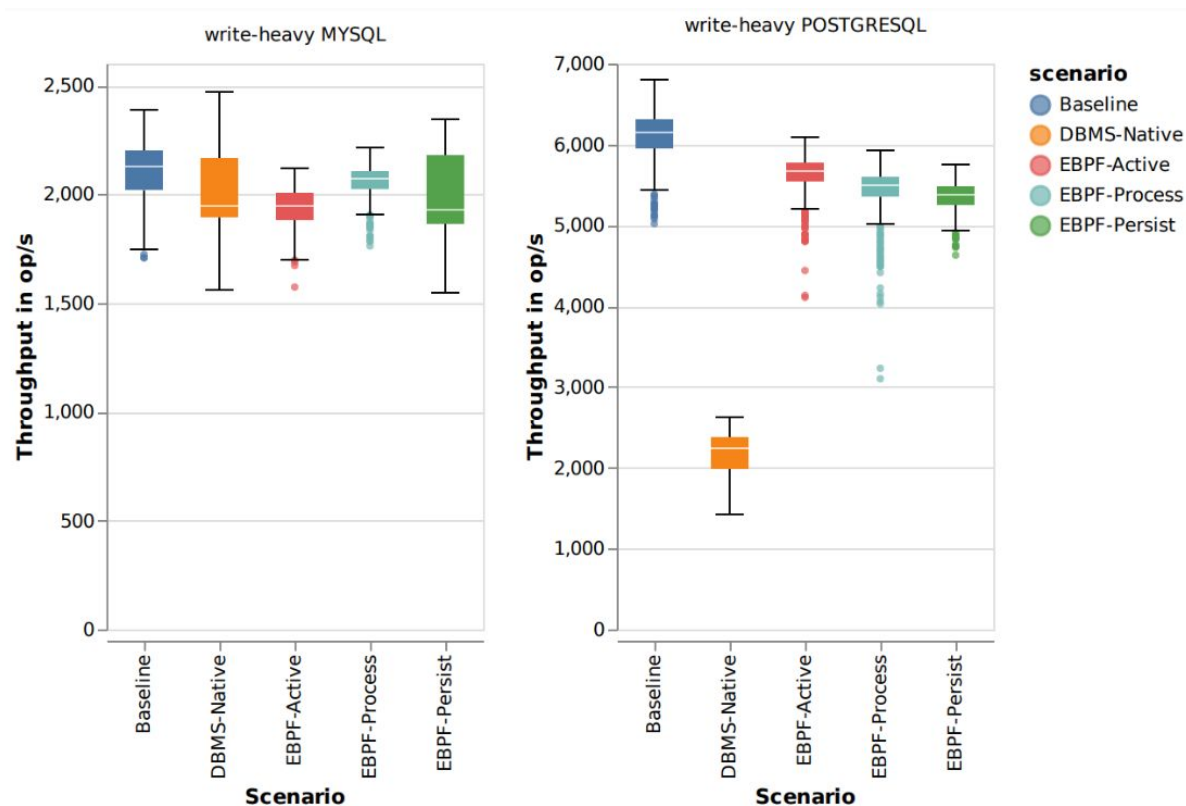
Process and persist queries to file with ANTtrail

## Evaluation Setup


	PostgreSQL	MySQL	YCSB
cloud	AWS EC2		
region	eu-central-1		
instance	m5.large	m5.large	c5.4xlarge
storage type	GP2		
OS	Ubuntu 20.04		
version	13.9	8.0.30	0.17.0

	read-heavy	write-heavy
YCSB instances	1	
threads	50	
inital data size	10 GB	
write proportion	0.1	0.9
read proportion	0.9	0.1
runtime	30 minutes	

## Results – Write-Heavy Workload



## Results – Read-Heavy Workload



## Result Evaluation

- ▶ **RQ1** Instrumentation with eBPF using uprobes works.  
No technical constraints to use with other DBMS.
- ▶ **RQ2** Performance impact with eBPF-based workload tracing is not stable across different workloads.  
Overhead depends on DBMS technology under test.
- ▶ **RQ3** eBPF-based approach competes with DBMS-native tracing or outperforms it.

## Summary

- ▶ eBPF has similar or lower impact on database performance than native DBMS tracing
- ▶ Impact highly depends on applied workload
- ▶ Very specific to database implementation
- ▶ DBMS independent traces via eBPF possible
- ▶ Generation of real world database traces helps in operations
- ▶ Our approach provides a non-intrusive method
- ▶ May be usable in production for selected cases

## Outlook


- ▶ Add more DBMS
  - ▶ Like MongoDB and Redis
- ▶ Improvements and optimizations of the eBPF program
- ▶ Take various probes into account e.g. kprobes
  - ▶ Retrieve query on signal path
- ▶ Different benchmarks and workloads
  - ▶ E.g. TPC

Thank you for attention!

Questions?



Source code and data available at:  
<https://github.com/benchANT/dbms-tracing-overhead>





# Results

**Table 4: Degradation for tracing a workload on MYSQL**

type variable scenario	read-heavy				write-heavy			
	$\Delta\%$ median throughput	median throughput	mean throughput	$\Delta\%$ mean throughput	$\Delta\%$ median throughput	median throughput	mean throughput	$\Delta\%$ mean throughput
Baseline	0.0	6432.00	6384.03	0.0	0.0	2127.40	2102.87	0.0
DBMS-Native	14.2	5518.80	5456.49	14.5	8.6	1945.50	1999.43	4.9
EBPF-Persist	3.6	6200.65	6070.62	4.9	9.4	1928.00	1988.55	5.4
EBPF-Process	2.5	6268.65	6162.64	3.5	2.7	2070.45	2059.28	2.1
EBPF-Active	9.2	5839.10	5921.38	7.2	8.5	1946.15	1936.57	7.9

# Results

**Table 5: Degradation for tracing a workload on POSTGRESQL**

type variable scenario	read-heavy				write-heavy			
	$\Delta\%$ median throughput	median throughput	mean throughput	$\Delta\%$ mean throughput	$\Delta\%$ median throughput	median throughput	mean throughput	$\Delta\%$ mean throughput
Baseline	0.0	11546.90	10076.10	0.0	0.0	6149.80	6109.23	0.0
DBMS-Native	21.6	9057.00	9080.74	9.9	63.5	2245.75	2180.93	64.3
EBPF-Persist	16.4	9656.45	9405.72	6.7	12.6	5374.05	5347.53	12.5
EBPF-Process	19.2	9331.80	9209.57	8.6	10.7	5490.00	5413.67	11.4
EBPF-Active	18.1	9460.30	9297.67	7.7	7.9	5666.45	5620.40	8.0